

Summary - Multi-Agent Distributed Lifelong Learning for Collective Knowledge Acquisition

Original File: 1709.05412v2.pdf | Generated on: June 12, 2026

1. Overall Summary

This paper introduces CoLLA, a novel algorithm for multi-agent distributed lifelong learning, addressing the limitations of existing single-agent, centralized lifelong learning approaches. CoLLA enables a network of agents to collectively learn a series of tasks by sharing learned knowledge in a distributed and decentralized manner, while preserving the privacy of locally observed data. The problem is formulated as an online multi-task learning optimization over a network, solved using an extended ADMM algorithm. The authors provide theoretical guarantees for robust performance and empirically demonstrate CoLLA's superiority over existing distributed multi-task learning methods on various datasets, including land mine detection, facial expression recognition, London schools, and computer surveys. Experimental results show that collaboration significantly improves both learning speed and asymptotic performance, with more connected network topologies leading to faster learning rates. The approach offers a practical solution for real-world scenarios where data is distributed, communication is constrained, or privacy is a concern.

2. Main Idea

The paper proposes Collective Lifelong Learning Algorithm (CoLLA), a distributed and decentralized multi-agent framework that enables a network of agents to collectively acquire knowledge over a series of tasks. The core idea is to allow agents to share learned knowledge efficiently and privately, improving overall task performance and learning speed, extending the concept of lifelong learning from a single agent to a collaborative multi-agent system.

3. Problem Being Solved

The paper addresses the limitations of current lifelong learning algorithms, which typically rely on a single learning agent with centralized access to all data. This centralized paradigm is impractical for many real-world applications where task data is distributed, partially accessible, requires local processing due to privacy concerns (e.g., healthcare), or involves costly/time-consuming data communication (e.g., home service robots, wearable devices). The paper aims to bridge this gap by enabling multiple agents to collectively learn a series of tasks in a distributed and decentralized fashion.

4. Assumptions

- The system consists of multiple collaborating lifelong learning agents.
- Each agent faces its own series of tasks, which may be unique to that agent.
- There is a latent knowledge base that underlies all tasks, and each agent learns a local view of this base.
- Agents are homogeneous, meaning there is no central node guiding collaboration.
- Agents are synchronous, meaning they simultaneously receive and learn one task at each time step.
- The communication mode between agents is modeled by a connected, undirected graph, imposing communication constraints.
- For theoretical convergence, the data distribution has a compact support, the LASSO problem in Eq. (4) admits a unique solution, and the matrices $L_i^T \Gamma_i(t) L_i$ are strictly positive definite (implying strong convexity of $F_i(t)$).

5. Limitations

- The paper explicitly states it does not address the privacy considerations that may arise from transferring learned knowledge between agents, although it preserves the privacy of locally observed raw data.
- The current framework assumes synchronous agents and a fixed, known order for agents, which is restrictive.
- The proposed approach is not yet extended to asynchronous agents with dynamic links, which is identified as future work.
- The computational cost of updating all local dictionaries by the agents is $O(Nu^2d^3)$, and the overall cost is $O(Nu^2(d, M) + K(Nu^2d^3 + eud))$, which could be high for very large N , u , d , or M , or many inner loop iterations K .

6. Results and Conclusion

The experimental results demonstrate that CoLLA significantly outperforms existing approaches like Single-Task Learning (STL), ELLA, and GO-MTL in multi-agent distributed lifelong learning settings across various classification and regression datasets. CoLLA consistently shows improved lifelong learning, both in terms of learning speed (faster convergence to a performance threshold) and asymptotic performance. The "jumpstart" analysis confirms CoLLA's effectiveness in transferring knowledge, leading to better initial performance on new tasks. Offline CoLLA's performance is comparable to the batch MTL method GO-MTL, and CoLLA achieves similar asymptotic performance to ELLA, indicating its effectiveness in both online and distributed settings. Furthermore, the study on graph structures reveals that more connected network topologies (e.g., complete graph) facilitate faster learning rates due to increased communication and collaboration among agents. The conclusion is that collaboration among agents, enabled by CoLLA, not only enhances asymptotic performance but also accelerates learning by requiring less data to achieve a specific performance threshold.

7. Real World Impact

This research has practical applications in scenarios requiring distributed intelligence and collaborative learning, such as: Financial Decision Support (agents access partial data views, collaboratively improving models); Smart Healthcare (local data processing due to privacy regulations); Autonomous Robotics/Edge AI (robots/wearable devices process perceptual data locally, sharing insights); Big Data Systems/Cloud Computing (parallel processing across multiple virtual agents for large-scale data analysis); and Personal Assistant Robots (agents collaboratively learn tasks like facial expression recognition).

8. Graphs and Figures Analysis

Figure 1: Performance of distributed (dotted lines), centralized (solid), and single-task learning (dashed) algorithms on benchmark datasets. Shaded region shows standard error (Best viewed in color.)

This figure presents four sub-plots (a-d) comparing the performance of CoLLA (proposed), ELLA (single-agent lifelong learning), Dist. CoLLA (offline distributed), GO-MTL (batch multi-task learning), and STL (single-task learning) across different benchmark datasets: Land Mine, Facial Expression, Computer Survey, and London Schools. The x-axis represents the number of learned tasks, and the y-axis represents the performance metric (AUC for classification in (a) and (b), RMSE for regression in (c) and (d)). Shaded regions indicate the standard error. In all plots, CoLLA (orange dotted line) consistently shows superior or comparable performance to other online methods (ELLA, STL) and approaches the performance of offline methods (Dist. CoLLA, GO-MTL) as the number of tasks increases. Specifically, CoLLA and Dist. CoLLA generally achieve higher AUC and lower RMSE, indicating better accuracy and lower error, respectively, compared to ELLA and STL. The performance of CoLLA and ELLA tends to converge asymptotically, as they solve the same underlying optimization problem. The figure highlights the benefits of collaboration and lifelong learning in improving performance over time.

Figure 2: Performance of CoLLA given various graph structures (a) for three data sets (b-d).

This figure investigates the impact of different network communication topologies on CoLLA's performance. Sub-figure (a) illustrates four graph structures: Random, Complete, Tree, and Server, which define how agents communicate. Sub-figures (b), (c), and (d) show CoLLA's performance (AUC for Facial Expression, RMSE for Computer Survey and London Schools, respectively) as a function of the number of tasks or AUC for these different graph structures. The results indicate that network structures with more edges, allowing for greater communication and collaboration, generally lead to faster learning rates and better performance. For instance, the "Complete" graph (where all agents are connected) often shows the best performance, followed by "Server" and "Random" graphs, while the "Tree" structure (minimal connectivity) tends to perform less optimally or converge slower. This suggests that increased communication among agents enhances the collective learning process.

Table 1: Jumpstart comparison (improvement in percentage) on the Land Mine (LM), London Schools (LS), Computer Survey(CS), and Facial Expression (FE) datasets

This table presents a "jumpstart" comparison, which quantifies the percentage improvement in initial performance on a new task due to knowledge transfer from previously learned tasks. It compares CoLLA, ELLA, Dist. CoLLA, and GO-MTL across four datasets: Land Mine (LM), London Schools (LS), Computer Survey (CS), and Facial Expression (FE). Higher percentages indicate better knowledge transfer and initial performance. CoLLA consistently shows strong jumpstart performance, often outperforming ELLA and sometimes even Dist. CoLLA and GO-MTL, particularly on datasets like Facial Expression and Computer Survey. This highlights CoLLA's effectiveness in collaboratively learning knowledge bases that are suitable for transfer, demonstrating that collaboration is most effective in learning initial tasks and leveraging past experience.