

Summary - Multi-Agent Distributed Lifelong Learning for Collective Knowledge Acquisition

Original File: 1709.05412v2.pdf | Generated on: June 13, 2026

1. Overall Summary

This paper introduces the Collective Lifelong Learning Algorithm (COLLA), a novel framework for multi-agent distributed lifelong learning. It addresses the limitations of traditional lifelong learning, which typically relies on a single agent with centralized data access, by enabling a network of agents to collectively learn a series of tasks in a distributed and decentralized manner. COLLAs allows agents to share learned knowledge, rather than raw data, thereby preserving local data privacy. The problem is formulated as an online Multi-Task Learning (MTL) optimization over a network, leveraging a parametric approach and extending existing single-agent online optimization techniques. The authors provide theoretical guarantees for COLLAs robust performance and empirically demonstrate its superiority over existing distributed multi-task learning approaches on various datasets, including land mine detection, facial expression recognition, London schools, and computer surveys. Experimental results show that collaboration among agents significantly improves both learning speed and asymptotic performance, with denser communication networks leading to faster learning. COLLAs also exhibits a strong "jumpstart" capability, indicating effective knowledge transfer to new tasks. This research contributes a scalable and privacy-preserving solution for collective knowledge acquisition in dynamic multi-agent environments.

2. Main Idea

The central concept is to extend lifelong learning from a single agent to a network of multiple agents that collectively acquire knowledge over a series of tasks. This is achieved through a distributed and decentralized algorithm (COLLA) that enables agents to share learned knowledge efficiently while preserving the privacy of their local data, thereby improving collective task performance and learning speed.

3. Problem Being Solved

This paper aims to solve the limitations of current lifelong learning algorithms, which typically rely on a single learning agent with centralized access to all data. This centralized paradigm is impractical for many real-world applications where multi-modal task data is partially accessible, local data processing is inevitable due to privacy regulations (e.g., healthcare), data communication is costly or time-consuming (e.g., home service robots, wearable devices), or parallel processing is essential for big data systems. The paper addresses the research gap of enabling multiple agents to collectively learn a series of tasks in a distributed and decentralized manner while preserving the privacy of locally observed data.

4. Assumptions

- Agents are synchronous, meaning they simultaneously receive and learn one task at each time step.
- There is a latent knowledge base that underlies all tasks, and each agent learns a local view of this knowledge base based on its own task distribution.
- The network graph representing communication among agents is connected.
- The data distribution has a compact support, which enforces boundedness on ridge estimators and Hessian matrices.
- The LASSO problem for sparse coefficients admits a unique solution.
- The matrices involved in the optimization are strictly positive definite, ensuring strong convexity of the objective functions.

5. Limitations

- The paper does not address privacy considerations that may arise from transferring knowledge between agents, focusing instead on preserving privacy of locally observed data.
- The arbitrary order of agents is restrictive and needs to be known and fixed.
- The current framework primarily focuses on collaborative agents that receive sequential tasks, despite potential extendability to parallel processing systems.
- Future work is needed to extend the proposed framework to asynchronous agents with dynamic links to improve real-world applicability.

6. Results and Conclusion

Experimental results demonstrate that COLLA consistently outperforms existing single-agent online lifelong learning (ELLA) and single-task learning (STL) algorithms across various classification and regression datasets. COLLA's performance is comparable to or even surpasses centralized offline multi-task learning (GO-MTL) and distributed offline multi-task learning (Dist. CoLLA) as the number of learned tasks increases. Collaboration among agents significantly improves both the learning speed and the asymptotic performance. The study also reveals that network structures with more edges (denser communication) lead to faster learning rates. Furthermore, COLLA exhibits a substantial "jumpstart" capability, indicating its effectiveness in transferring learned knowledge to improve initial performance on new tasks. The final conclusion is that COLLA provides an efficient, robust, distributed, and decentralized approach for multi-agent lifelong learning, enabling collective knowledge acquisition while preserving local data privacy.

7. Real World Impact

This research has practical applications in various domains, including financial decision support systems where agents have partial data access, healthcare systems requiring local data processing due to privacy regulations, home service robots that need to process perceptual data locally, wearable devices with limited communication bandwidth, and big data systems that necessitate parallel processing across multiple virtual agents. It also applies to personal assistant robots for tasks like facial expression recognition and general autonomous agents and multiagent systems requiring collective knowledge acquisition in dynamic, uncertain environments.

8. Graphs and Figures Analysis

Figure 1: Performance of distributed (dotted lines), centralized (solid), and single-task learning (dashed) algorithms on benchmark datasets. Shaded region shows standard error (Best viewed in color.)

This figure presents the performance comparison of various learning algorithms (COLLA, ELLA, Dist. CoLLA, GO-MTL, STL) on four benchmark datasets: Land Mine, Facial Expression, Computer Survey, and London Schools. For classification tasks (Land Mine, Facial Expression), performance is measured by AUC (Area Under the ROC Curve), where higher is better. For regression tasks (Computer Survey, London Schools), performance is measured by RMSE (Root Mean Squared Error), where lower is better. The x-axis represents the number of learned tasks. COLLA (distributed online) consistently shows superior performance compared to ELLA (single-agent online) and STL (single-task learning), with its performance improving as more tasks are learned. Dist. CoLLA (distributed offline) and GO-MTL (centralized offline) generally serve as upper bounds, and COLLA often approaches or matches their performance with increasing tasks. The shaded regions indicate the standard error over 100 trials, demonstrating the robustness of the results. Overall, the figure illustrates that collaborative multi-agent lifelong learning (COLLA) effectively leverages knowledge transfer to achieve better and more stable performance than single-agent or single-task approaches.

Figure 2: Performance of CoLLA given various graph structures (a) for three data sets (b-d).

This figure investigates the impact of different network graph structures on the performance of the CoLLA algorithm across three datasets: Facial Expression, Computer Survey, and London Schools. Figure 2a visually depicts the four graph structures considered: Random, Complete, Tree, and Server (star graph). Figures 2b-2d show the learning performance (AUC for Facial Expression, RMSE for Computer Survey and London Schools) as a function of the number of tasks for each graph structure. The results consistently indicate that denser communication networks, such as 'Complete' (all agents connected) and 'Server' (agents connected to a central master agent), lead to faster learning rates and better overall performance (higher AUC or lower RMSE). Sparser networks like 'Tree' and 'Random' generally exhibit slower learning and slightly inferior performance. This suggests that increased communication and collaboration among agents, facilitated by more connected graph topologies, significantly enhance the learning speed and effectiveness of the distributed lifelong learning process.

Table 1: Jumpstart comparison (improvement in percentage) on the Land Mine (LM), London Schools (LS), Computer Survey(CS), and Facial Expression (FE) datasets

This table presents a 'jumpstart' comparison, which quantifies the percentage improvement in initial performance on a new task due to knowledge transfer from previously learned tasks. The comparison is made across four datasets (Land Mine, London Schools, Computer Survey, Facial Expression) for COLLA, ELLA, Dist. CoLLA, and GO-MTL. COLLA consistently demonstrates significant positive jumpstart percentages across all datasets (6.87% for Land Mine, 29.62% for London Schools, 51.44% for Computer Survey, and 40.87% for Facial Expression). ELLA also shows positive jumpstart but generally to a lesser extent than COLLA. Dist. CoLLA and GO-MTL, being offline methods, show very high jumpstart values, acting as strong upper bounds. The results highlight COLLA's effectiveness in collaboratively learning knowledge bases that are suitable for transfer, enabling agents to quickly adapt and perform well on new tasks by leveraging collective past experience.