

Summary - Multi-Agent Distributed Lifelong Learning for Collective Knowledge Acquisition

Original File: 1709.05412v2.pdf | Generated on: June 15, 2026

1. Overall Summary

This paper introduces the Collective Lifelong Learning Algorithm (COLLA), a novel approach for multi-agent distributed lifelong learning. It addresses the limitations of traditional lifelong learning, which typically relies on a single agent with centralized data access. COLLA enables a network of agents to collaboratively acquire knowledge from a series of tasks in a distributed and decentralized manner, ensuring local data privacy. The core idea is that knowledge learned by one agent can benefit others learning related tasks. The authors formulate this as an online multi-task learning optimization problem over a network, using a parametric approach and solving it with an extended Alternating Direction Method of Multipliers (ADMM) algorithm. This allows agents to share learned knowledge bases with neighbors without sharing raw data or specific model parameters. Theoretical guarantees for robust performance are provided, and empirical evaluations on various datasets demonstrate that COLLA outperforms existing distributed multi-task learning methods, showing improved learning speed and asymptotic performance, especially in initial task learning (jumpstart). The research highlights the benefits of collaboration among agents in dynamic, real-world environments and suggests future work on asynchronous agents and dynamic links.

2. Main Idea

The central concept is to extend lifelong learning from a single agent to a network of multiple agents that collectively acquire knowledge from a series of tasks in a distributed and decentralized manner, while preserving the privacy of locally observed data. The paper proposes the Collective Lifelong Learning Algorithm (COLLA) to achieve this.

3. Problem Being Solved

This paper aims to solve the weakness of previous lifelong learning systems that rely on a single learning agent with centralized access to all data. This centralized paradigm is impractical for many real-world applications where data is inherently distributed, partially accessible, or subject to privacy regulations that prevent sharing. The research gap addressed is the lack of a robust, distributed, and decentralized multi-agent lifelong learning framework that allows agents to collectively learn and share knowledge efficiently while preserving local data privacy.

4. Assumptions

- Agents in the network are synchronous.
- There is a latent knowledge base that underlies all tasks, and each agent learns a local view of this knowledge base.
- The network graph connecting agents is connected.
- All nodes (agents) in the network are homogeneous.
- The data distribution for tasks has a compact support, ensuring boundedness of parameters.
- The LASSO problem used for sparse coefficient estimation admits a unique solution.
- The matrices involved in the optimization ($L(t)$) are strictly positive definite, implying strong convexity of the loss functions.

5. Limitations

- The paper does not address privacy considerations that may arise from transferring knowledge between agents.
- The current focus is on collaborative agents that receive sequential tasks, not parallel processing systems, despite potential extendability.
- The arbitrary order on the agents in the network is restrictive and needs to be known and fixed for the current framework.
- The proposed framework is limited to synchronous agents.
- The framework does not yet extend to a network of asynchronous agents with dynamic links, which is identified as a future direction.
- The inner loop for agreement (K iterations) needs to be sufficiently large, especially initially, which implies a communication cost.

6. Results and Conclusion

Experimental findings demonstrate that COLLA consistently outperforms existing distributed multi-task learning approaches, including single-task learning (STL), single-agent lifelong learning (ELLA), and even distributed offline multi-task learning (Dist. CoLLA) and centralized offline multi-task learning (GO-MTL) in certain aspects. Collaboration among agents significantly improves lifelong learning, leading to faster learning speeds and better asymptotic performance compared to single-agent methods. COLLA exhibits a strong 'jumpstart' effect, indicating efficient knowledge transfer to new tasks, particularly benefiting initial task learning. The performance of offline CoLLA is comparable to GO-MTL, validating the algorithm's effectiveness in a distributed batch setting. Furthermore, network structures with more edges facilitate faster learning rates, suggesting that increased communication and collaboration among agents accelerate knowledge acquisition. The inner loop for agreement can be decreased over time as agents' local dictionaries converge.

7. Real World Impact

This research has practical applications in developing autonomous agents such as personal intelligent robot assistants and chatbots that can continuously interact with users and adapt to dynamic environments. It can be applied to financial decision support agents with partial data access, healthcare systems where personal medical data cannot be shared, home service robots processing large volumes of perceptual data locally, and wearable devices with limited communication bandwidth. Furthermore, it is relevant for modern big data systems that necessitate parallel processing in the cloud across multiple virtual agents (CPUs or GPUs) and general collective knowledge acquisition in various societal and biological contexts.

8. Graphs and Figures Analysis

Figure 1: Performance of distributed (dotted lines), centralized (solid), and single-task learning (dashed) algorithms on benchmark datasets. Shaded region shows standard error (Best viewed in color.)

This figure presents four sub-plots (a-d) comparing the performance of different learning algorithms (COLLA, ELLA, Dist. CoLLA, GO-MTL, STL) across four benchmark datasets: Land Mine (a), Facial Expression (b), Computer Survey (c), and London Schools (d). The x-axis represents the 'Number of tasks' learned by a single agent, and the y-axis represents the performance metric (AUC for classification tasks like Land Mine and Facial Expression, RMSE for regression tasks like Computer Survey and London Schools). COLLA, the proposed online distributed lifelong learning algorithm, consistently shows a strong learning curve, improving performance as more tasks are learned. It generally outperforms ELLA (single-agent lifelong learning without communication) and STL (single-task learning). ELLA shows improvement over STL but is typically outperformed by COLLA, especially in the initial learning phases, highlighting the benefit of multi-agent collaboration. Dist. CoLLA (offline distributed lifelong learning) and GO-MTL (offline centralized multi-task learning) serve as upper-bounds, generally achieving the highest performance. STL consistently shows the lowest performance. For all datasets, performance (AUC or inverse RMSE) generally increases with the number of tasks, demonstrating positive knowledge transfer. COLLA's performance quickly approaches that of the offline methods (Dist. CoLLA, GO-MTL) as more tasks are learned, indicating effective online distributed learning. The shaded regions show the standard error, suggesting the results are robust.

Figure 2: Performance of CoLLA given various graph structures (a) for three data sets (b-d).

This figure consists of four sub-plots. Sub-plot (a) visualizes four different network graph structures: Random, Complete, Tree, and Server (star graph), which define the communication topology among agents. Sub-plots (b), (c), and (d) show the performance of COLLA on the Facial Expression, Computer Survey, and London Schools datasets, respectively, under these different graph structures. The x-axis represents the 'Number of tasks' (#AUC or #task), and the y-axis represents the performance metric (AUC for Facial Expression, RMSE for Computer Survey and London Schools). For all three datasets, COLLA's performance generally improves with the number of tasks learned, regardless of the graph structure. However, graph structures with more edges (e.g., Complete, Server) consistently lead to faster learning rates and often higher asymptotic performance compared to sparser structures (e.g., Tree, Random). This indicates that increased communication and collaboration among agents, facilitated by denser network topologies, accelerates knowledge acquisition and improves collective learning efficiency. The Server and Complete graph structures often show the best performance, while the Tree structure, representing minimal connectivity, typically shows slower learning.

Table 1: Jumpstart comparison (improvement in percentage) on the Land Mine (LM), London Schools (LS), Computer Survey(CS), and Facial Expression (FE) datasets

This table presents a 'jumpstart' comparison, which quantifies the initial performance improvement on a new task due to knowledge transfer, expressed as a percentage. It compares COLLA, ELLA, Dist. CoLLA, and GO-MTL across four datasets: Land Mine (LM), London Schools (LS), Computer Survey (CS), and Facial Expression (FE). COLLA demonstrates significant jumpstart improvements across all datasets (e.g., 51.44% on CS, 40.87% on FE), highlighting its effectiveness in transferring learned knowledge to new tasks. ELLA, while also showing improvements, generally achieves slightly less jumpstart than COLLA, particularly on CS and FE, underscoring the benefit of multi-agent collaboration in COLLA. Dist. CoLLA (offline distributed) and GO-MTL (offline centralized) achieve the highest jumpstart improvements, serving as strong baselines for batch learning. The table clearly demonstrates that COLLA effectively leverages collaborative learning to achieve strong 'jumpstart' performance, indicating efficient knowledge transfer to new tasks, and providing an advantage over single-agent lifelong learning (ELLA) in initial task performance.