

Summary - Multi-Agent Distributed Lifelong Learning for Collective Knowledge Acquisition

Original File: 1709.05412v2.pdf | Generated on: June 13, 2026

1. Overall Summary

This paper introduces CoLLA (Collective Lifelong Learning Algorithm), a novel distributed and decentralized optimization algorithm for multi-agent lifelong learning. It addresses the limitations of traditional lifelong learning, which relies on a single agent with centralized data access, by enabling a network of agents to collectively learn a series of tasks. The problem is formulated as an online Multi-Task Learning (MTL) optimization over a network, where agents share learned knowledge while keeping their local data private. CoLLA extends existing single-agent lifelong learning techniques by incorporating a second-order Taylor approximation and an online optimization scheme within a distributed framework. The algorithm provides theoretical guarantees for robust performance and convergence. Empirical evaluations on diverse datasets, including land mine detection, facial expression recognition, London schools, and computer survey, demonstrate that CoLLA consistently outperforms existing distributed multi-task learning approaches. The results show that collaboration significantly improves both learning speed and asymptotic performance, with more connected network topologies leading to faster learning. CoLLA's offline performance is comparable to state-of-the-art batch MTL methods, highlighting its effectiveness in both online and offline settings and its ability to facilitate knowledge transfer for initial tasks.

2. Main Idea

The central concept is to extend lifelong learning from a single agent to a network of multiple agents that collectively acquire knowledge over a series of tasks in a distributed and decentralized manner, while preserving the privacy of locally observed data. The paper proposes CoLLA, a Collective Lifelong Learning Algorithm, to achieve this by enabling efficient knowledge sharing among agents.

3. Problem Being Solved

This paper addresses the limitations of previous lifelong learning systems that rely on a single learning agent with centralized access to all data. It tackles the research gap concerning scenarios where multi-modal task data is only partially accessible by individual agents, local data processing is mandated by privacy regulations (e.g., healthcare), data communication is costly or time-consuming, or parallel processing is essential due to large data sizes or geographical distribution. The paper aims to solve the challenge of enabling multiple agents to collectively learn a series of distributed and potentially unique tasks, sharing learned knowledge in a decentralized and privacy-preserving manner.

4. Assumptions

- The hidden mapping from each data point to its corresponding target can be parametrized.
- A linear mapping is considered, though the framework is generalizable to nonlinear parametric mappings.
- There is a latent knowledge base underlying all tasks, and each agent learns a local view of this base.
- Each agent solves a local version of the objective to estimate its own local knowledge base.
- Agents are synchronous, meaning they simultaneously receive and learn one task at each time step.
- The network graph is connected to allow knowledge and information flow between all agents.
- All nodes (agents) are homogeneous, with no central node guiding collaboration.
- The data distribution has a compact support, ensuring boundedness of certain parameters.
- The LASSO problem in Equation (4) admits a unique solution, making the functions $F_i(t)$ well-defined.
- The matrices $L_i^T * \Gamma_i(t) * L_i$ are strictly positive definite, implying the functions $F_i(t)$ are strongly convex.

5. Limitations

- The paper does not address privacy considerations that may arise from transferring knowledge between agents, focusing instead on preserving privacy of locally observed data.
- The arbitrary order of agents is restrictive and needs to be known and fixed.
- The current framework focuses on collaborative agents that receive sequential tasks, despite potential extendability to parallel processing systems.
- The proposed framework is currently limited to synchronous agents with static links, with asynchronous agents and dynamic links being a direction for future work.

6. Results and Conclusion

The experimental findings demonstrate that CoLLA consistently outperforms existing single-agent lifelong learning (ELLA) and single-task learning (STL) algorithms across various classification and regression datasets (Land Mine, Facial Expression, London Schools, Computer Survey). Collaboration among agents significantly improves both the learning speed and asymptotic performance. The offline version of CoLLA achieves performance comparable to the GO-MTL batch multi-task learning algorithm, indicating its effectiveness as a distributed MTL approach. CoLLA also exhibits a 'jumpstart' effect, showing strong initial performance gains on new tasks due to effective knowledge transfer. Furthermore, the study reveals that network structures with more edges lead to faster learning rates, implying that increased communication and collaboration enhance learning speed. The algorithm is supported by theoretical guarantees for robust performance and convergence to a stationary point of the risk function.

7. Real World Impact

This research has practical applications in various domains where multiple agents need to collaborate and learn from distributed data. Examples include financial decision support systems where agents have access to only partial data views, healthcare systems requiring local data processing due to privacy regulations, home service robots processing perceptual data locally due to volume, and wearable devices with limited communication bandwidth. It can also be applied in modern big data systems that necessitate parallel processing across multiple virtual agents and in personal assistant robots for tasks like facial expression recognition.

8. Graphs and Figures Analysis

Figure 1: Performance of distributed (dotted lines), centralized (solid), and single-task learning (dashed) algorithms on benchmark datasets.

This figure presents the performance of various learning algorithms (CoLLA, ELLA, Dist. CoLLA, GO-MTL, STL) on four benchmark datasets as the number of learned tasks increases. The shaded regions represent the standard error. For classification tasks (Land Mine and Facial Expression), higher AUC (Area Under the ROC Curve) indicates better performance. For regression tasks (Computer Survey and London Schools), lower RMSE (Root Mean Squared Error) indicates better performance. In all sub-figures, CoLLA (online, distributed) consistently outperforms ELLA (single-agent lifelong learning) and STL (single-task learning), demonstrating the benefits of multi-agent collaboration and lifelong learning. CoLLA's performance approaches that of Dist. CoLLA (offline, distributed) and GO-MTL (batch multi-task learning), which serve as upper bounds, as more tasks are learned. This indicates that CoLLA effectively leverages shared knowledge in a distributed online setting.

Figure 2: Performance of CoLLA given various graph structures (a) for three data sets (b-d).

This figure illustrates the impact of different network graph structures on CoLLA's performance across three datasets. Figure 2a visualizes the four graph structures: Random, Complete, Tree, and Server. Figures 2b-2d show that network structures with more edges, such as 'Complete' and 'Server', generally lead to faster learning rates and better overall performance (higher AUC for Facial Expression, lower RMSE for Computer Survey and London Schools) compared to less connected structures like 'Tree' and 'Random'. This suggests that increased communication and collaboration among agents, facilitated by denser network topologies, significantly enhance the learning speed and effectiveness of the CoLLA algorithm.