

Summary - Multi-Agent Distributed Lifelong Learning for Collective Knowledge Acquisition

Original File: 1709.05412v2.pdf | Generated on: June 13, 2026

1. Overall Summary

This paper introduces the Collective Lifelong Learning Algorithm (CoLLA), a novel approach for multi-agent distributed lifelong learning. It addresses the limitation of traditional lifelong learning, which typically relies on a single agent with centralized data access, by enabling a network of agents to collectively learn a series of tasks. CoLLA facilitates efficient, distributed, and decentralized knowledge sharing among agents, allowing them to leverage knowledge from related tasks to improve individual and collective learning performance, all while preserving the privacy of locally observed data. The authors formulate the learning problem as an online Multi-Task Learning (MTL) optimization over a network, where each agent learns parametric models for its tasks and shares its knowledge base with neighbors. The Alternating Direction Method of Multipliers (ADMM) algorithm is adapted to solve this global optimization problem in an online, distributed manner. Theoretical guarantees for robust performance and convergence are provided. Empirical evaluations on various benchmark datasets (Land Mine, Facial Expression, Computer Survey, London Schools) demonstrate that CoLLA consistently outperforms existing distributed multi-task learning and single-agent lifelong learning approaches. The results show that collaboration significantly improves both learning speed and asymptotic performance, and that denser network topologies lead to faster learning. CoLLA also exhibits strong 'jumpstart' performance, indicating effective knowledge transfer to new tasks. The research highlights the potential for scalable and adaptive learning in dynamic, multi-agent environments.

2. Main Idea

The central concept is to extend lifelong learning from a single agent to a network of multiple agents that collectively acquire knowledge. This is achieved through a distributed and decentralized algorithm (CoLLA) that allows agents to share learned knowledge efficiently while preserving the privacy of local data, thereby improving task performance and learning speed across a series of tasks.

3. Problem Being Solved

This paper addresses the limitations of previous lifelong learning systems that rely on a single learning agent with centralized access to all data. It tackles the research gap where many real-world applications involve scenarios with multiple agents, distributed tasks, and constraints such as partial data accessibility, privacy regulations preventing data sharing, costly or time-consuming data communication, and the necessity for parallel processing due to large data sizes or geographical distribution. The paper aims to solve the challenge of enabling these multiple agents to collectively learn a series of tasks in a distributed and decentralized manner, sharing knowledge effectively while maintaining data privacy.

4. Assumptions

- Agents are synchronous, meaning they simultaneously receive and learn one task at each time step.
- The network graph representing communication channels between agents is connected.
- All nodes (agents) in the network are homogeneous.
- There is a latent knowledge base that underlies all tasks, and each agent learns a local view of this knowledge base.
- The data distribution has a compact support, which enforces boundedness on ridge estimators and Hessian matrices, and subsequently on local knowledge bases and sparse coefficients.
- The LASSO problem in Equation (4) admits a unique solution.
- The matrices $L_i^T * \Gamma(t) * L_i$ are strictly positive definite, ensuring that the functions $F_i(t)$ are all strongly convex.
- The order of agents is arbitrary but needs to be known and fixed.

5. Limitations

- The proposed approach does not address the privacy considerations that may arise from transferring knowledge between agents.
- The current focus is on collaborative agents that receive sequential tasks, despite potential extendability to parallel processing systems.
- The framework assumes synchronous agents; extending it to a network of asynchronous agents with dynamic links is a potential future direction.
- The arbitrary order of agents is restrictive and needs to be known and fixed.

6. Results and Conclusion

The experimental findings demonstrate that the Collective Lifelong Learning Algorithm (CoLLA) consistently outperforms existing distributed multi-task learning and single-agent lifelong learning approaches across various benchmark datasets (Land Mine, Facial Expression, Computer Survey, London Schools). Collaboration among agents significantly improves lifelong learning, leading to both faster learning speeds and better asymptotic performance. CoLLA's performance is comparable to offline, centralized methods like GO-MTL when used in a distributed batch MTL setting. Furthermore, CoLLA shows superior 'jumpstart' performance, indicating its effectiveness in transferring knowledge to new tasks. The research also concludes that network structures with more edges facilitate faster learning rates. The algorithm is theoretically proven to converge to a stable dictionary and a stationary point of the risk function, validating its robust performance.

7. Real World Impact

This research has practical real-world applications in various domains, including financial decision support systems where agents may have partial data access, smart healthcare where regulations prevent sharing personal medical data, home service robots that need to process perceptual data locally due to volume, wearable devices with limited communication bandwidth, and modern big data systems requiring parallel processing in the cloud across multiple virtual agents. It also applies to personal intelligent robot assistants and chatbots that continuously interact with different users, enabling them to adapt and improve performance over time in dynamic environments.

8. Graphs and Figures Analysis

Figure 1: Performance of distributed (dotted lines), centralized (solid), and single-task learning (dashed) algorithms on benchmark datasets.

This figure presents the performance of different lifelong learning algorithms across four benchmark datasets: Land Mine (a), Facial Expression (b), Computer Survey (c), and London Schools (d). The y-axis represents AUC (Area Under the ROC Curve) for classification tasks (a, b) and RMSE (Root Mean Squared Error) for regression tasks (c, d), while the x-axis denotes the number of tasks learned. For AUC, higher is better; for RMSE, lower is better. The shaded regions indicate the standard error over 100 trials. CoLLA (Collective Lifelong Learning Algorithm, dotted blue circles) and Dist. CoLLA (Distributed CoLLA, dotted green stars) consistently show superior performance compared to ELLA (Efficient Lifelong Learning Algorithm, solid red squares), GO-MTL (Group-Online Multi-Task Learning, solid purple diamonds), and STL (Single-Task Learning, dashed black triangles). Specifically, CoLLA and Dist. CoLLA achieve higher AUC or lower RMSE faster and reach better asymptotic performance, demonstrating the benefits of multi-agent collaboration and distributed learning. Dist. CoLLA and GO-MTL, being offline methods, generally achieve the best asymptotic performance, while online CoLLA is competitive and outperforms other online approaches like ELLA.

Figure 2: Performance of CoLLA given various graph structures (a) for three data sets (b-d).

This figure investigates the impact of different network graph structures on CoLLA's performance. Figure 2a visualizes four graph structures: Random, Complete, Tree, and Server. Figures 2b-2d show CoLLA's performance (AUC for Facial Expression, RMSE for Computer Survey and London Schools) as a function of the number of tasks for these different graph structures. The results indicate that network structures with more edges, such as 'Complete' and 'Server' topologies, lead to faster learning rates and improved performance (higher AUC or lower RMSE) compared to sparser structures like 'Tree' and 'Random'. This suggests that increased communication and collaboration among agents, facilitated by denser network connectivity, enhance the learning speed and overall effectiveness of the CoLLA algorithm.

Table 1: Jumpstart comparison (improvement in percentage) on the Land Mine (LM), London Schools (LS), Computer Survey(CS), and Facial Expression (FE) datasets

This table presents a 'jumpstart' comparison, showing the percentage improvement in initial performance on a new task due to knowledge transfer, across four datasets: Land Mine (LM), London Schools (LS), Computer Survey (CS), and Facial Expression (FE). The algorithms compared are CoLLA, ELLA, Dist. CoLLA, and GO-MTL. CoLLA consistently demonstrates significant jumpstart improvements (e.g., 6.87% for LM, 29.62% for LS, 51.44% for CS, 40.87% for FE), indicating its effectiveness in leveraging collaboratively learned knowledge bases for new tasks. As expected, offline methods like Dist. CoLLA and GO-MTL generally show higher jumpstart values. Crucially, CoLLA's jumpstart performance is superior to ELLA, highlighting the advantage of collaborative learning in accelerating initial task performance.